

Introduction to Computation for the Humanities and Social Sciences



CS 3

Chris Tanner

Lecture 16

Nothing says autumn like
some fresh ~~apple pie~~
numpy



Lecture 16

- Numpy

- Pandas

Numerical Analysis and Computation on Large Datasets

- Python is incredibly powerful and essentially comes with functionality to do anything we could want
- However, when doing computations with a lot of numerical data, there's an external library which provides us with even more power (faster operations and more convenient/succinct code):
Numpy and **PANDAS**!

Numpy

Numpy (numb-pie)

- Python lists do not have the notion of addition, subtraction, multiplying, and dividing lists of numbers. We'd have to manually write these functions and iterate through every number.
- Numpy provides this functionality, and it can compute such *very fast*
- Numpy was designed to facilitate computing on large amounts of **num**erical data in **py**thon.

REAL-TIME CODING

Numpy examples

Numpy

Numpy

- Numpy is an external library, so we must import it if we want to use it.

```
import numpy as np
```

- We didn't have to call it "**np**" but that's common practice. Whatever name we gave that variable, that's what we use to access the numpy functions

Numpy — Data Structures

- Python's `list` data structure allows us to store a list of data. Numpy provides such too and calls their data structure an `array`.
- The basic data structures in numpy are `array` and `matrices`

Numpy — Array

- You can construct a numpy `array` from a list:

```
import numpy as np
weights = [121, 130, 220, 170, 148]

ages = np.array([21, 18, 19, 20, 19])
new_weights = np.array(weights)
```

Numpy — Array

- numpy `array`'s do not support different data types in the same array; if you try to do such, numpy will guess a type and convert the entire array to that type

```
import numpy as np
```

```
# the array below will be casted to contain strings  
np.array([1, 2.5, 'hello'])
```

```
# the array below will be casted to contain floats  
np.array([1, 2.5])
```

Numpy — Array

- We can do many operations with `array`'s:
 - `np.min()`
 - `np.max()`
 - `np.sum()`
 - `np.mean()`
 - `np.median()`
 - `np.var()`
 - `np.std()`
 - `np.sin()` or `np.cos()` or `np.tan()`
 - etc

Numpy — Array

- `np.arange()` is similar to the range function we've been using
- except it allows us to use floats for the start, stop, and step.
- and, we can see all of the values ahead of time:

```
np.arange(10)
```

```
np.arange(0, 1, 0.1)
```

Numpy — Matrices

- Matrix is a two dimensional array of data, a la a list of lists, but array of arrays.
- All arrays it contains must be of the same length.
- Can create a matrix of all zeros via:

```
np.zeros ( num_rows , num_columns )
```

Numpy — Matrices

- Indexing into a matrix is similar to an array, but we add now have to specify the row and column
- Now we have `m[row, column]`
- It supports the same start, stop, step syntax we used with a python list and a numpy array, just that the row and column are separated by a comma

`m[0:10, 0]`

Numpy — Matrices

- Will get the first column of data. We can shorten this with just a single colon
- Similarly, to get the first row of data, we can switch them:

`m[0, :]`

Numpy — Matrices

- One of the cool things that numpy allows us to do is assign a value or a list to a slice of values. For example, if we wanted the first row of the matrix to be 1s we can say:

```
m[0, :] = 1
```


Numpy — Resources

- For more info on Numpy, check out their website:
<https://docs.scipy.org/doc/numpy/user/quickstart.html>

Lecture 16

- Numpy

- Pandas

Motivation

pandas



Motivation

pandas

- **pandas** is a data analysis library built on top of **numpy** and many other libraries.
- It extends Numpy in cool and interesting ways, and it comprises your homework

Motivation

pandas

- As with numpy, the first thing we must do is import pandas. Again we use the as syntax to shorten our commands.
- You'll also often find examples online referring just to commands with pd or np. That implies that pandas and numpy were imported in this manner.

```
import pandas as pd
```

pandas — Series

- The pandas Series is similar to the numpy array in that it's a one-dimensional data array
- It even supports the Numpy functions, like mean, sum, max etc
- Creating a pandas Series is simply

```
series = pd.Series(np.arange(10))
```

pandas — Series

- pandas supports giving a customized index label to each value

```
series = pd.Series(np.arange(5), index=[ 'zero', 'one', 'two', 'three', 'four' ])
```

- Now we can refer to the value in the series either by its numerical index or its label index

```
series[0]    or  
series[ 'zero' ]
```

pandas — Series

- Can even make a series from a Python dictionary

```
dict = {'b' : 1, 'a' : 0, 'c' : 2}  
series = pd.series(dict)
```


pandas — DataFrame

- The **DataFrame** is a more extensible version of the Numpy 2-d matrix. It allows columns to be different types, which the numpy matrix didn't.
- This makes it very convenient to read DataFrame directly from a **.csv** file!

```
data_frame = pd.read_csv("data.csv")
```

pandas — DataFrame

- Can take a quick look at examples of the data by calling:

```
data_frame.head( )
```

- You can get a quick summary of the data by calling:

```
data_frame.describe( )
```

pandas — DataFrame

- If our .csv file includes a header row of labels, then we can refer to the columns by their header label!

```
data_frame[ 'happiness' ]
```

- To refer to specific rows, we can address a range of them:

```
data_frame[ 0 : 3 ]
```

pandas — DataFrame

- To get a specific row, we can use the following
(which creates a very readable look at the row of data)

```
data_frame.iloc[0]
```

- Can address a single data point:

```
data_frame.at[0, 'Generosity']
```

- You can get the column labels by looking at

```
data_frame.columns
```

Motivation

pandas — DataFrame

Unlike numpy data structures, you can add columns to pandas data frames like you would a Python dictionary

```
data_frame['Unhappiness Score'] = 10.0 - data_frame['Happiness Score']
```

pandas — DataFrame

Like a numpy array or matrix, you can use a filter syntax

```
data_frame[data_frame['happiness'] > 7]
```

Then if you wanted just the list of countries:

```
data_frame[data_frame['happiness'] > 7]['Country']
```

pandas — DataFrame

Summarizing operations by default only occur on a single axis

```
data_frame.mean()
```

If you want to summarize by the other axis

```
data_frame.max(1)
```

Motivation

pandas — DataFrame

After creating a data frame, you can save it to a csv file with simply
`data_frame.to_csv('output.csv')`

pandas even supports reading and writing to excel files!

`df.to_excel('foo.xlsx', sheet_name='Sheet1')`

`pd.read_excel('foo.xlsx', 'Sheet1', index_col=None, na_values=['NA'])`

LAB TIME

